

# class 12

Faisal

#Import Data

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG00000000003	723	486	904	445	1170
ENSG00000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2

	SRR1039517	SRR1039520	SRR1039521
ENSG00000000003	1097	806	604
ENSG00000000005	0	0	0
ENSG000000000419	781	417	509
ENSG000000000457	447	330	324
ENSG000000000460	94	102	74
ENSG000000000938	0	0	0

Q1 How many genes are in this dataset?

```
nrow(counts)
```

[1] 38694

Q2 How many 'control' cell lines do we have?

```
ncol(counts)
```

```
[1] 8
```

and the metadata aka “colData”

```
(metadata)
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867
5	SRR1039516	control	N080611	GSM1275870
6	SRR1039517	treated	N080611	GSM1275871
7	SRR1039520	control	N061011	GSM1275874
8	SRR1039521	treated	N061011	GSM1275875

Lets make sure that the id column of the metadata match the order of the columns in Count-Data.

```
metadata$id == colnames(counts)
```

```
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

We can use the 'all()' function to check that all its input are TRUE

```
all( c(T,T,T, F))
```

```
[1] FALSE
```

```
all( metadata$id == colnames(counts))
```

```
[1] TRUE
```

## Analysis by hand

```
metadata
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867
5	SRR1039516	control	N080611	GSM1275870
6	SRR1039517	treated	N080611	GSM1275871
7	SRR1039520	control	N061011	GSM1275874
8	SRR1039521	treated	N061011	GSM1275875

Lets first extract our counts for control samples to compare this to the count for treated (i.e with drug) samples

Q3. How would you make the above code in either approach more robust?

```
control.inds <- metadata$dex == "control"
control.ids <- metadata$id[ control.inds]
control.counts <- counts[, control.ids ]
control.mean <- rowMeans(control.counts)
head(control.counts)
```

	SRR1039508	SRR1039512	SRR1039516	SRR1039520
ENSG000000000003	723	904	1170	806
ENSG000000000005	0	0	0	0
ENSG000000000419	467	616	582	417
ENSG000000000457	347	364	318	330
ENSG000000000460	96	73	118	102
ENSG000000000938	0	1	2	0

I want a single summary counts value for each gene in the control experiments. I will start by taking the average

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean)

```
##apply(control.counts, 1, mean)
treated.mean <- rowMeans(control.counts)
treated.inds <- metadata$dex == "treated"
treated.ids <- metadata$id[ control.inds]
treated.counts = counts[, treated.ids ]
head(treated.counts)
```

	SRR1039508	SRR1039512	SRR1039516	SRR1039520
ENSG00000000003	723	904	1170	806
ENSG00000000005	0	0	0	0
ENSG00000000419	467	616	582	417
ENSG00000000457	347	364	318	330
ENSG00000000460	96	73	118	102
ENSG00000000938	0	1	2	0

```
treated.mean = rowMeans(treated.counts)
```

Now we do the same for the treated samples Please :-)

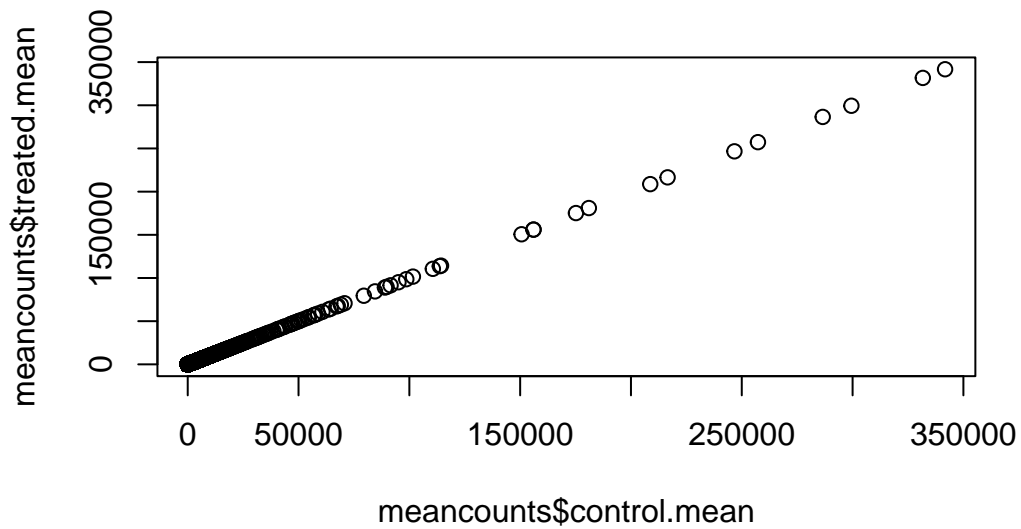
```
meancounts <- data.frame(control.mean, treated.mean)
head(meancounts)
```

	control.mean	treated.mean
ENSG00000000003	900.75	900.75
ENSG00000000005	0.00	0.00
ENSG00000000419	520.50	520.50
ENSG00000000457	339.75	339.75
ENSG00000000460	97.25	97.25
ENSG00000000938	0.75	0.75

and make a wee plot to see how we are doing

Q5 Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.

```
plot(meancounts$control.mean, meancounts$treated.mean)
```



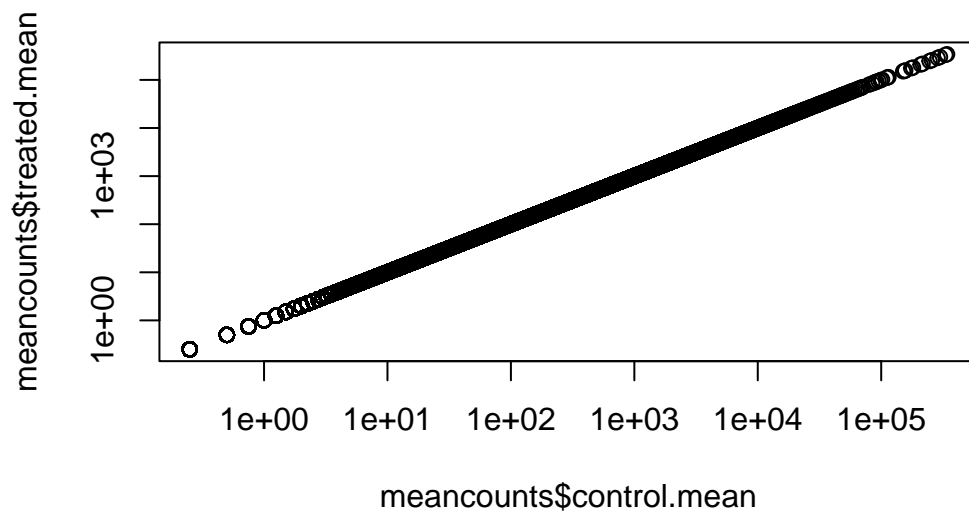
This screams for a log transformation so we can see our data

Q6 Try plotting both axes on a log scale. What is the argument to `plot()` that allows you to do this?

```
plot(meancounts$control.mean, meancounts$treated.mean, log="xy")
```

Warning in `xy.coords(x, y, xlabel, ylabel, log)`: 15032 x values  $\leq 0$  omitted from logarithmic plot

Warning in `xy.coords(x, y, xlabel, ylabel, log)`: 15032 y values  $\leq 0$  omitted from logarithmic plot



The most useful and most straightforward to understand is log2 transformation

```
log2(20/20)
```

```
[1] 0
```

Doubling

```
log2(40/20)
```

```
[1] 1
```

```
log2(10/20)
```

```
[1] -1
```

add a “log2 fold-change”

```
meancounts$log2fc <- log2(meancounts$treated.mean / meancounts$control.mean)
```

```
head(meancounts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	900.75	0
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	520.50	0
ENSG000000000457	339.75	339.75	0
ENSG000000000460	97.25	97.25	0
ENSG000000000938	0.75	0.75	0

Hmmm... we need to get rid of the genes where we have no count data as taking the log2 of these 0 counts does not tell us anything.

```
head( meancounts == 0)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	FALSE	FALSE	TRUE
ENSG000000000005	TRUE	TRUE	NA
ENSG000000000419	FALSE	FALSE	TRUE
ENSG000000000457	FALSE	FALSE	TRUE
ENSG000000000460	FALSE	FALSE	TRUE
ENSG000000000938	FALSE	FALSE	TRUE

```
to.keep <- rowSums(meancounts[,1:2] == 0) == 0
```

```
mycounts <- meancounts[to.keep,]
head(mycounts)
```

	control.mean	treated.mean	log2fc
ENSG000000000003	900.75	900.75	0
ENSG000000000419	520.50	520.50	0
ENSG000000000457	339.75	339.75	0
ENSG000000000460	97.25	97.25	0
ENSG000000000938	0.75	0.75	0
ENSG000000000971	5219.00	5219.00	0

Q7. What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call the `unique()` function?

it returns the true value.. we use the unique value because we don't need the position that has 2 trees, which is repeated.

How many genes are up regulated at the `log2fc` level of +2

Q8. Using the `up.ind` vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

```
sum(mycounts$log2fc >= +2)
```

```
[1] 0
```

Q9. Using the `down.ind` vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

and down regulated...

```
sum(mycounts$log2fc <= -2)
```

```
[1] 0
```

Q10 No we dont trust these results because we dont know if the numbers are significant

We are missing the stats..

## DESeq2 analysis

```
library(DESeq2)
```

Like most bioconductor packages DESeq wants its input and output in a very specific format

```
dds <- DESeqDataSetFromMatrix(countData=counts,  
                              colData=metadata,  
                              design=~dex)
```



converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

```
dds
```

```
class: DESeqDataSet
dim: 38694 8
metadata(1): version
assays(1): counts
rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
               ENSG00000283123
rowData names(0):
colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
colData names(4): id dex celltype geo_id
```

The main DESeq function is called DESeq

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control


DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG000000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029
	padj				
	<numeric>				
ENSG000000000003	0.163035				
ENSG000000000005	NA				
ENSG000000000419	0.176032				
ENSG000000000457	0.961694				
ENSG000000000460	0.815849				
ENSG000000000938	NA				

## Volcano plots


A major summary figure of this type of analysis is called a volcano plot - the idea here is to keep our inner biologist and inner stats person happy with one cool plot

```
plot( res$log2FoldChange, res$padj)
```




Improve this plot by taking the log of that p-value axis

```
plot( res$log2FoldChange,  log(res$padj) )
```



I want to flip this y-axis so that the value i care about are at the top of the axis

```
plot( res$log2FoldChange,  -log(res$padj),  
      xlab="Log2(FoldChange)",  
      ylab="-Log(P-value)")
```



## gene annotation

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
columns(org.Hs.eg.db)
```

```
[1] "ACCNUM"      "ALIAS"       "ENSEMBL"     "ENSEMBLPROT" "ENSEMBLTRANS"
[6] "ENTREZID"    "ENZYME"      "EVIDENCE"    "EVIDENCEALL"  "GENENAME"
[11] "GENETYPE"    "GO"          "GOALL"       "IPI"          "MAP"
[16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL" "PATH"         "PFAM"
[21] "PMID"        "PROSITE"     "REFSEQ"      "SYMBOL"       "UCSCKG"
[26] "UNIPROT"
```

```
#Pathway anaylysis
```

```
library(pathview)
```

```
#####  
Pathview is an open source software package distributed under GNU General  
Public License version 3 (GPLv3). Details of GPLv3 is available at  
http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to  
formally cite the original Pathview paper (not just mention it) in publications  
or products. For details, do citation("pathview") within R.
```

The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG  
license agreement (details at <http://www.kegg.jp/kegg/legal.html>).

```
#####
```

```
library(gage)
```

```
library(gageData)
```

```
data("kegg.sets.hs")
```

```
#examine the first 2 pathways in this kegg set for human
```

```
head(kegg.sets.hs, 2)
```

```
$`hsa00232 Caffeine metabolism`
```

```
[1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
$`hsa00983 Drug metabolism - other enzymes`
```

```
[1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"  
[9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"  
[17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"  
[25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"  
[33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"  
[41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"  
[49] "8824" "8833" "9" "978"
```

```
c(barry=4, clair=3, chandra=2)
```

```
barry    clair chandra
      4      3      2
```

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
[1] -0.35070302          NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

```
keggres = gage(foldchanges, gsets=kegg.sets.hs)
attributes(keggres)
```

```
$names
[1] "greater" "less"    "stats"
```

```
head(keggres$less, 3)
```

		p.geomean	stat.mean	p.val	q.val
hsa00232	Caffeine metabolism	NA	NaN	NA	NA
hsa00983	Drug metabolism - other enzymes	NA	NaN	NA	NA
hsa01100	Metabolic pathways	NA	NaN	NA	NA

		set.size	expl
hsa00232	Caffeine metabolism	0	NA
hsa00983	Drug metabolism - other enzymes	0	NA
hsa01100	Metabolic pathways	0	NA

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

Warning: None of the genes or compounds mapped to the pathway!  
Argument gene.idtype or cpd.idtype may be wrong.

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/faissal/Desktop/UCSD/BIMM 143/Class 15

Info: Writing image file hsa05310.pathview.png

I put this in the document

